

OPERATING SYSTEMS

Lecture # 6

Razi Uddin



INTER-PROCESS COMMUNICATION (IPC)

IPC provides a mechanism to allow processes to communicate and to synchronize their actions without sharing the same address space.



MESSAGE PASSING SYSTEM

- The function of a message system is to allow processes to communicate without the need to resort to shared data.
- Messages sent by a process may be of either fixed or variable size.
- If processes P and Q want to communicate, a communication link must exist between them and they must send messages to and receive messages from each other through this link.



MESSAGE PASSING SYSTEM

Methods for logically implementing a link and the send and receive options:

- Direct or indirect communication
- Symmetric or asymmetric communication
- Automatic or explicit buffering
- Send by copy or send by reference
- Fixed-size or variable size messages



DIRECT COMMUNICATION

With direct communication, each process that wants to communicate must explicitly name the recipient or sender of the communication.

- The send and receive primitives are defined as:
- `Send(P, message)` – send a message to process P
- `Receive(Q, message)` – receive a message from process Q.



DIRECT COMMUNICATION

A communication link in this scheme has the following properties:

- A link is established automatically between every pair of processes that want to communicate.
- The processes need to know only each other's identity to communicate
- A link is associated with exactly two processes.
- Exactly one link exists between each pair of processes.



DIRECT COMMUNICATION

Unlike this symmetric addressing scheme, a variant of this scheme employs asymmetric addressing, in which the recipient is not required to name the sender.

- `Send(P, message)` – send a message to process P
- `Receive(id, message)` – receive a message from any process; the variable `id` is set to the name of the process with which communication has taken place.



INDIRECT COMMUNICATION

- With indirect communication, messages can be sent to and received from mailboxes.
- Here, two processes can communicate only if they share a mailbox.
- The send and receive primitives are defined as:
 - `Send(A, message)` – send a message to mailbox A.
 - `Receive(A, message)` – receive a message from mailbox A.



INDIRECT COMMUNICATI ON

- A communication link in this scheme has the following properties:
- A link is established between a pair of processes only if both members have a shared mailbox.
- A link is associated with more than two processes.
- A number of different links may exist between each pair of communicating processes, with each link corresponding to one mailbox.



SYNCHRONIZATION

- Communication between processes takes place by calls to send and receive primitives (i.e., functions).
- Message passing may be either blocking or non-blocking also called synchronous and asynchronous.

Blocking send—The sending process is blocked until the receiving process or the mailbox receives the message.

Non-blocking send—The sending process sends the message and resumes operation.

Blocking receiver—The receiver blocks until a message is available.

Non-blocking receiver—The receiver receives either a valid message or a null.



BUFFERING

- Whether the communication is direct or indirect, messages exchanged by the processes reside in a temporary queue. This queue can be implemented in three ways:
- **Zero Capacity:** The queue has a maximum length zero, thus the link cannot have any messages waiting in it. In this case, the sender must block until the message has been received.
- **Bounded Capacity:** This queue has finite length n ; thus at most n messages can reside in it. If the queue is not full when a new message is sent, the latter is placed in the queue and the sender resumes operation. If the queue is full, the sender blocks until space is available.
- **Unbounded Capacity:** The queue has infinite length; thus the sender never blocks.



UNIX/LINUX IPC TOOLS

UNIX and Linux operating systems provide many tools for inter-process communication.

- Pipe
- Named pipe (FIFO)
- BSD Socket
- TLI
- Message queue
- Shared memory
- Etc.



OPEN() SYSTEM CALL

- The open() system call is used to open or create a file. Its synopsis is as follows:

```
#include<sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
```

```
int open(const char pathname, int oflag, /* mode_t mode */);
```



OPEN() SYSTEM CALL

- The call converts a pathname into a file descriptor (a small, non-negative integer for use in subsequent I/O as with read, write, etc.).
- When the call is successful, the file descriptor returned will be the lowest file descriptor not currently open for the process.
- This system call can also specify whether read or write will be blocking or non-blocking.
- The 'oflag' argument specifies the purpose of opening the file and 'mode' specifies permission on the file if it is to be created.
- 'oflag' value is constructed by ORing various flags: O_RDONLY, O_WRONLY, O_RDWR, O_NDELAY (or O_NONBLOCK), O_APPEND, O_CREAT, etc



READ SYSTEM CALL

`read()`—attempts to read up to `count` bytes from file descriptor `fd` into the buffer starting at `buf`. If `count` is zero, `read()` returns zero and has no other results. If `count` is greater than `SSIZE_MAX`, the result is unspecified. On success, `read()` returns the number of bytes read (zero indicates end of file) and advances the file position pointer by this number.

`#include<unistd.h>`

`ssize_t read(int fd, void *buf, size_t count);`



WRITE SYSTEM CALL

`write()`—attempts to write up to `count` bytes to the file referenced by the file descriptor `fd` from the buffer starting at `buf`. On success, `write()` returns the number of bytes written (zero indicates nothing was written) and advances the file position pointer by this number. On error, `write()` returns `-1`, and `errno` is set appropriately. If `count` is zero and the file descriptor refers to a regular file, 0 will be returned without causing any other effect.

`#include <unistd.h>`

`ssize_t write(int fd, const void *buf, size_t count);`



CLOSE SYSTEM CALL

`close()`—closes a file descriptor, so that it no longer refers to any file and may be reused. If `fd` is the last copy of a particular file descriptor the resources associated with it are freed; if the descriptor was the last reference to a file which has been removed using `unlink(2)` the file is deleted. `close()` returns zero on success, or `-1` if an error occurred.

```
#include <unistd.h>
```

```
int close(int fd);
```




PIPES

- A UNIX/Linux pipe can be used for IPC between related processes on a system.
- Communicating processes typically have sibling or parent-child relationship.
- At the command line, a pipe can be used to connect the standard output of one process to the standard input of another.
- Pipes provide a method of one-way communication and for this reason may be called half-duplex pipes.
- The `pipe()` system call creates a pipe and returns two file descriptors, one for reading and second for writing.
- The files associated with these file descriptors are streams and are both opened for reading and writing.
- Naturally, to use such a channel properly, one needs to form some kind of protocol in which data is sent over the pipe.
- Also, if we want a two-way communication, we'll need two pipes.



PIPES



PIPES

- The system assures us of one thing: the order in which data is written to the pipe, is the same order as that in which data is read from the pipe.
- The system also assures that data won't get lost in the middle, unless one of the processes (the sender or the receiver) exits prematurely.
- The `pipe()` system call is used to create a read-write pipe that may later be used to communicate with a process we'll fork off.
- The synopsis of the system call is:

`#include<unistd.h>`

`int pipe (int fd[2]);`




PIPES

- `fd[0]` is the file descriptor for the read end of the pipe (i.e., the descriptor to be used with the `read` system call),
- whereas `fd[1]` is the file descriptor for the write end of the pipe. (i.e., the descriptor to be used with the `write` system call).
- The function returns `-1` if the call fails.
- A pipe is a bounded buffer and the maximum data written is `PIPE_BUF`, defined in `<sys/param.h>` in UNIX and in `<linux/param.h>` in Linux as 5120 and 4096, respectively.



PIPES



Half-duplex pipe after a fork.

